RIVR – Texas Hold'em
Iteration Planning, Execution and Testing

CS 482, Team 23
Chris, Finn, Kelechi and Royal

November 3, 2026

Purpose and provenance. This document repairs the missing Iteration 1 documentation and records Iteration 2 work. Iteration 1 is reconstructed retrospectively; it is not represented as a report submitted before the original deadline. Team plans were supplied by Chris. Branch commits are contribution evidence; main-branch delivery and professor acceptance are separate checks. Individual reflections, actual hours and acceptance decisions require human confirmation.

The Iteration 2 assignment supplied in chat lists November 5, 2026 as the nominal due date, with instructions to verify Moodle and a Sunday-night grace period. The exact cutoff, release date and presentation time must be verified; the demo is not a delivery extension. Planning-email content has been supplied, but its sent timestamp/receipt has not.

Contents

1 System Design **2**

2 Main Iteration Planning Table **4**

3 Iteration 1: Planning and Execution **4**

4 Iteration 2: Planning and Execution **5**

5 Testing and Coverage Report **6**

6 Process, Traceability and Retrospective **9**

7 Release and Demo Checklist **9**

1 System Design

RIVR follows the existing three-layer design: React/TypeScript presentation, Django/REST Framework application services, and PostgreSQL persistence. Session authentication and CSRF protect account writes. Server-owned event IDs, transactions and user-row locks protect XP and chip rewards. Avatar bytes are validated and normalized before storage.

1.1 Diagrams

The following sequence diagram models a proposed multiplayer action and hand-completion scenario. It shows invalid-action rollback, authoritative result calculation, transactional XP/trophies and recipient-filtered publication only after commit. The game service, game-action endpoint and Channels transport are planned, not implemented by this modelling deliverable. The progression service and underlying framework/database exist now. A diagram is not evidence of a shipped gameplay engine.

Editable source: `docs/diagrams/multiplayer-sequence.mmd`. Vector rendering is generated by `scripts/render_gameplay_diagram.py`. The model deliberately distinguishes shared pots from outright wins and keeps private hole cards private.

RIVR: proposed multiplayer action and hand completion

Design model only. Game endpoints and Channels are not implemented; progression hooks exist.



Invalid actions never award XP. Shared pots are not wins. Broadcast only after commit; private cards remain private.

2 Main Iteration Planning Table

Iteration	Owner	Planned stories	Points
1	Chris	S13 public table; S10 browse tables	4
1	Finn	S2 registration; S1 login/logout; S4 guests	4
1	Kelechi	S23 admin cleanup; S35 themes	4
1	Royal	S5 image upload; S14 starting chips; S29 daily chips	4
1	Team	Updated team plan	16
2	Chris	S32 XP/achievements (3); S18 leaderboard (1)	4
2	Finn	S19 add/remove friends	5
2	Kelechi	S16 chip balance (2); S15 buy-in (1); S17 stats (1)	4
2	Royal	S11 filters (1); S27 private table (1); S26 spectate (2)	4
2	Team	Reported planning-email content	17

These are planned points, not accepted points. The original workbook lists S5 as 3 points and S14 as unestimated; Chris supplied an updated Iteration 1 team plan of 2 and 1. Instructor estimate confirmation is not supplied. Planned ownership is distinct from commit authorship. Some S11 filtering and chip-persistence behavior existed earlier; the team must identify incremental work and avoid counting the same delivery twice.

3 Iteration 1: Planning and Execution

3.1 Reported plan

The original planning window was October 12–22, 2026, with an October 21 internal completion target and an October 22 nominal deadline to confirm on Moodle. Four points were assigned to each teammate, for 16 planned points. The report preserves the updated team plan rather than inventing a prior submission date. Planning-email timing remains *Pending actual records.*

3.2 Reconstructed work

Stories	Planned owner	Observed source / original authorship evidence
S13/S10	Chris	Tables, create/lobby screens and tests; S13 commit `9b87b7d`.
S2/S1/S4	Finn	Authentication/session and frontend commits `402f931`, `03cb335`, `a8606f1`.
S23/S35	Kelechi	Cleanup/theme commits `54bae1f`, `d3e56f2`, `f0ac4cb`.
S5/S14/S29	Royal assigned	Chris later covered these as requested in this session. `9b21c1e` is authored by Chris on October 26. It is not attributed to Royal or backdated.

All ten feature areas exist in the current code, but this does not prove their presence at the original cutoff. The fetched main tip at the start of this audit was `7a3d567`, October 26. The October 26 avatar/reward work cannot be claimed as a demonstrated original-deadline delivery. Other teammates' local/unpushed work cannot be established from this clone alone.

Actual hours for each story, instructor acceptance, total accepted points, velocity and historical cycle time: *Pending actual records.* The original 14–16 hours/person estimate is not an actual time

log. The CSV work-log template preserves unknowns as blank, not zero.

3.3 Testing and reflection

The professor reported that the committed coverage image covered only two files. That historical S23 report covers cleanup/admin only (61 statements and 12 branches), not the team application. It is retained with an explicit scope correction. Current full-application evidence is supplied separately and does not retroactively cure the original submission.

Proposed process improvements, awaiting individual team confirmation: write the report during the iteration; check every file included in coverage; link real issues in commits; rehearse the real database flow; preserve authorship when merging; and ensure every teammate attends and explains their own work. The professor reported Kelechi presented alone. Attendance explanations belong to the students and must not be manufactured.

4 Iteration 2: Planning and Execution

4.1 Team branches and contributions

Chris confirmed each member works on their own branch. Fetched branch refs and original commit authors are recorded in `docs/quality/branch-evidence.json`. Finn's `feature/S19-friends` includes `798af8e`, `09166e7` and `e3c7733`; it was not in the audited main tip. Kelechi's historical cleanup/theme branches and Chris's `saia` are visible. A missing or unpushed Royal/current teammate ref is not evidence of non-contribution. The team must provide those refs and integrate reviewed work by the cutoff.

4.2 Chris: S32 and S18

S32, 3 planned points, issue #45. Twenty-four trophies have hidden-until-earned metadata, private account collection access and server-owned rewards. Real avatar upload, public game creation and recorded theme changes award persisted XP. Each trophy unlocks once; Against the Odds pays 50 XP again per distinct qualifying 7–2 offsuit win. Outright hand wins pay 100 XP and tournament wins 1,000 XP per unique result. Awards stack and retries do not pay twice. Splits are not wins.

Gameplay, tournament, friends and messaging trophies have tested trusted hooks and explicit validation rules, but their authoritative feature integrations are not shipped by this work. This is a partial integration boundary, not proof of a functioning poker engine. The professor decides whether the demonstrated scope satisfies the story; three accepted points are not assumed.

S18, 1 planned point, issue #32. The all-time leaderboard aggregates actual ledger XP in PostgreSQL, excludes inactive users, includes zero-XP active accounts and uses shared competition ranks (1, 1, 3). ID ordering stabilizes tied pagination. The public payload contains username, avatar, ID, rank and XP only. Signed-in players can see their own global rank off-page. Inputs, page sizes, error states, refresh and privacy are tested. Friends/weekly/season rankings, levels and win rates are not fabricated.

4.3 Execution record and remaining team inputs

Owner	Current evidence / boundary	Actual hours
Chris	S32 core and three real action integrations; S18 implemented on saia; acceptance pending	Pending
Finn	S19 feature branch observed; team integration/acceptance pending	Pending
Kelechi	S16/S15/S17 assigned; teammate branch/results to confirm	Pending
Royal	S11/S27/S26 assigned; teammate branch/results to confirm	Pending

Team planned points are 17. Accepted/delivered point totals and individual actual hours remain pending. No extra points are claimed for optional refactoring, spikes, card animation or documentation repair. Branch activity is evidence of work; the rubric still requires the reviewed deliverable on main before the verified cutoff.

5 Testing and Coverage Report

Measured source commit: `06063bc902d8eae7ec3c4e6f21221bbfd054e49b`.

Measurement time (UTC): `2026-11-03T20:28:33.902615+00:00`.

Scope	Files	Statements	Branches
Backend	37	880/899 (97.89%)	201/208 (96.63%)
Frontend	33	541/561 (96.43%)	453/489 (92.63%)

JUnit evidence: backend 362 tests; frontend 115 tests. Failures/errors: 0 on each side.

Equal-weight means: statements 97.16%; branches 94.63%.

The manifest and raw reports provide exact denominators, missing code and exclusions. The two percentages are reported separately because Python executable-line statements and JavaScript statements have different denominators. The optional equal-weight means below are a transparent interpretation of the rubric, not a pooled cross-language percentage. Each side independently exceeds the target where the gate passes.

Backend measurement includes all Python under backend, including configuration and launchers. Tests, generated migrations and the virtual environment are excluded. Frontend includes executable TS/TSX under src, including main.tsx; tests, setup and declaration files are excluded. Type-only modules may compile to no executable code. Unexecuted runtime files remain in the reports. Supporting report-generation tools are additionally smoke-tested by generating and validating their artifacts; they are not silently counted as runtime app coverage.

The measurement describes this review branch, not unmerged teammate work. Regenerate after main integration. Frontend exports use relative paths only; counts and maps are unchanged. JUnit XML preserves actual test totals.

Raw JSON/XML, per-file images, source SHA/hashes, missing code and commands are committed under `docs/quality/evidence/`. The historical two-file image is not reused as team coverage. HTML reports are regenerated locally or uploaded by CI. Hosted CI success is not asserted until a pushed run is observed.

5.1 Tool-generated report prints

```
Backend full-application coverage
Measured runtime files (including unexecuted files)
Source commit: 06063bc902d8
```

File	Statements	Branches
accounts/__init__.py	0/0	0/0
accounts/admin.py	4/4	0/0
accounts/models.py	11/11	0/0
accounts/serializers.py	43/43	8/8
accounts/services.py	24/24	6/6
accounts/urls.py	4/4	0/0
accounts/views.py	68/68	6/6
config/__init__.py	0/0	0/0
config/asgi.py	0/4	0/0
config/settings.py	31/32	1/2
config/urls.py	10/10	1/2
config/wsgi.py	0/4	0/0
conftest.py	35/35	0/0
core/__init__.py	0/0	0/0
core/views.py	15/15	0/0
manage.py	0/8	0/2
profiles/__init__.py	0/0	0/0
profiles/avatars.py	59/59	9/10
profiles/rewards.py	42/42	10/10
profiles/urls.py	4/4	0/0
profiles/views.py	52/52	6/6
progression/__init__.py	0/0	0/0
progression/catalog.py	5/5	0/0
progression/facts.py	60/60	50/50
progression/leaderboard.py	50/50	10/10
progression/models.py	22/22	0/0
progression/services.py	71/73	40/42
progression/views.py	18/18	0/0
tables/__init__.py	0/0	0/0
tables/admin.py	26/26	4/4
tables/cleanup.py	35/35	8/8
tables/models.py	48/48	12/12
tables/pagination.py	11/11	0/0
tables/permissions.py	8/8	0/0
tables/serializers.py	58/58	14/14
tables/urls.py	4/4	0/0
tables/views.py	62/62	16/16
TOTAL	97.89%	96.63%

```
Tests, generated migrations and type-only declarations are excluded.
Backend includes config and launchers; frontend includes main.tsx.
See manifest.json and raw tool JSON for denominators and missing code.
```

Figure 1: Full backend per-file coverage generated from the tool's raw JSON.

```
Frontend full-application coverage
Measured runtime files (including unexecuted files)
Source commit: 06063bc902d8
```

File	Statements	Branches
src/App.tsx	23/24	13/16
src/api/client.ts	11/11	4/4
src/components/AppLayout.tsx	10/10	41/41
src/components/FormField.tsx	2/2	10/10
src/features/achievements/AchievementsPage.tsx	27/27	19/20
src/features/auth/AuthLayout.tsx	10/10	12/12
src/features/auth/GuestPage.tsx	9/9	2/2
src/features/auth/RegisterPage.tsx	14/14	6/6
src/features/auth/SessionGate.tsx	11/11	12/12
src/features/auth/SignInPage.tsx	9/9	2/2
src/features/auth/api.ts	12/12	0/0
src/features/auth/types.ts	0/0	0/0
src/features/auth/useAuthForm.ts	46/48	19/22
src/features/auth/validation.ts	29/29	24/24
src/features/leaderboard/LeaderboardPage.tsx	28/30	35/37
src/features/profile/AccountCard.tsx	34/41	31/41
src/features/profile/ProfileGate.tsx	8/8	10/10
src/features/profile/ProfileProvider.tsx	24/24	25/27
src/features/profile/RewardsPage.tsx	31/34	28/32
src/features/profile/api.ts	8/8	0/0
src/features/profile/context.ts	3/3	0/0
src/features/profile/types.ts	0/0	0/0
src/features/tables/CreateTablePage.tsx	52/52	48/52
src/features/tables/LobbyPage.tsx	61/63	58/62
src/features/tables/TableDetailPage.tsx	14/14	7/9
src/features/tables/api.ts	8/8	0/0
src/features/tables/types.ts	0/0	0/0
src/features/tables/validation.ts	20/20	28/28
src/features/themes/SettingsPage.tsx	18/18	16/16
src/features/themes/ThemeProvider.tsx	10/10	0/0
src/features/themes/theme.ts	5/5	2/2
src/features/themes/useTheme.ts	4/5	1/2
src/main.tsx	0/2	0/0
TOTAL	96.43%	92.63%

```
Tests, generated migrations and type-only declarations are excluded.
Backend includes config and launchers; frontend includes main.tsx.
See manifest.json and raw tool JSON for denominators and missing code.
```

Figure 2: Full frontend per-file coverage generated from the tool's raw JSON.

5.2 Adequacy, gaps and improvement

The test suite checks persistence, validation, CSRF, permissions, owner isolation, atomic failure/rollback, image content and cleanup, exact UTC claim boundaries, concurrent payouts, event retries and ranking ties/page boundaries. Frontend tests cover loading, retry, stale responses, invalid inputs and theme changes. Real-browser checks exercise actual registration, table/avatar/theme XP and ledger ranking in both themes and mobile/desktop layouts.

High coverage is useful regression evidence, not acceptance or feature completeness. Missing game, tournament and social integrations still need real end-to-end tests once their authoritative services exist. CLI bootstrap code may remain uncovered. Read the missing-line/branch lists rather than increasing a percentage through trivial mirrored tests. CI repeats full-scope checks and fails if either side does not exceed 80% statement and 50% branch coverage; reports remain truthful on failure.

6 Process, Traceability and Retrospective

The live Kanban Status options are Backlog, Ready, In progress, In review, Acceptance Test and Done. Acceptance Test was added immediately before Done while preserving all 39 existing statuses. Only professor/CTO acceptance authorizes Done. New commits reference actual issues #45 and #32; older hashes are mapped without rewriting original authorship.

Velocity is the accepted story-point total during the iteration. Cycle time uses actual In Progress and Acceptance Test timestamps under the team's convention. Neither metric can be reconstructed reliably from commit timestamps or newly corrected statuses; unknowns remain pending. The work log captures future actual records. Main delivery, peer review, submission timing and demo attendance are human/team responsibilities distinct from a local test run.

Proposed next-iteration improvements: verify the coverage scope before release; keep the planning/execution report current; record time and transitions when they occur; review small issue-linked commits; rehearse happy and failure paths with real data; disclose partial integrations early; and coordinate a backup presenter/recording without replacing required live attendance. Each teammate should add their own reflection rather than accepting generated prose as a personal statement.

7 Release and Demo Checklist

1. Confirm planning email timing, actual hours, professor acceptance history and exact Moodle/demo dates.
2. Push original feature branches; peer-review and integrate without losing authorship. Main must contain the release before cutoff.
3. Run the full quality script, export truthful evidence for the source SHA, build this report and commit the PDF/source/images.
4. Rehearse S32/S18 real actions and failure cases. Show the proposed gameplay model separately from running features.
5. Every member presents their own work and can explain scope, tests and commits. Communicate a real absence in advance.
6. Record CTO decisions after the demo; accepted stories may move to Done and rejected work returns to Backlog.